

# Stanford CS230 Notes

## 1 Introduction to Deep Learning

### 1.1 What is a Neural Network?

neuron, links.

### 1.2 Supervised Learning with Neural Networks

Supervised Learning

Examples: Standard NN, Convolutional NN, Recurrent NN

Structured Data: tabular data; Unstructured Data: audio/image/text

### 1.3 Why is Deep Learning taking off?

Amount of labeled data.

## 2 Basics of Neural Network Programming

### 2.1 Binary Classification

#### 2.1.1 Binary Classification

image  $\rightarrow$  1 (cat) vs 0 (non cat)

#### 2.1.2 Notation

$m$ : number of examples

$n_x$ : input size

$n_y$ : output size

$x$ : input, column vector

$y$ : output, 0/1

$X$ : input matrix, shape =  $(n_x, m)$

$Y$ : output matrix, shape =  $(1, m)$

$x^{(i)}$ : superscript (i) will denote the  $i^{th}$  example.

### 2.2 Logistic Regression

Given:  $x \in \mathbb{R}^{n_x}$ ,  $0 \leq \hat{y} \leq 1$

Parameters:  $w \in \mathbb{R}^{n_x}$ ,  $b \in \mathbb{R}$

Output:

$$z = w^T x + b \tag{1}$$

$$\hat{y} = \sigma(z) \tag{2}$$

$$\sigma(z) \approx \frac{1}{1 + e^{-z}} \tag{3}$$

$$z \approx \infty, \sigma(z) \approx \frac{1}{1 + 0} = 1 \tag{4}$$

$$z \approx -\infty, \sigma(z) \approx \frac{1}{1 + \infty} = 0 \tag{5}$$

Simplified Parameters:  $x_0 = 1$ ,  $x \in \mathbb{R}^{n_x+1}$

$$\theta_0 = b, \theta_1 \dots \theta_{n_x} = w \quad (6)$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \dots \\ \theta_{n_x} \end{bmatrix} \quad (7)$$

$$\hat{y} = \sigma(\theta^T x) \quad (8)$$

### 2.3 Logistic Regression cost function

Given  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$ , want  $\hat{y}^{(i)} \approx y^{(i)}$ .

Loss (error) function (mean square error/cross entropy):

$$\mathcal{L}(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2 \quad (9)$$

$$\mathcal{L}(\hat{y}, y) = -(y \log \hat{y} + (1 - y) \log(1 - \hat{y})) \quad (10)$$

If  $y = 1$ :  $\mathcal{L}(\hat{y}, y) = -\log \hat{y}$ , want  $\mathcal{L}$  small, want  $\hat{y}$  large, want  $\hat{y}$  equal to 1. If  $y = 0$ :  $\mathcal{L}(\hat{y}, y) = -\log(1 - \hat{y})$ , want  $\mathcal{L}$  small, want  $\hat{y}$  small, want  $\hat{y}$  equal to 0.

Cost function:

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}, y) \quad (11)$$

$$J(w, b) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] \quad (12)$$

Given a random variable  $X$  with probability mass function  $p_X(x)$ , the self information of measuring  $X$  as outcome  $x$  is defined as:

$$I_X(x) = \log[p_X(x)] = \log\left(\frac{1}{p_X(x)}\right) \quad (13)$$

Shannon Entropy of  $X$ :

$$H(X) = \sum_x -p_X(x) \log p_X(x) \quad (14)$$

$$= \sum_x p_X(x) I_X(x) \quad (15)$$

$$= E[I_X(x)] \quad (16)$$

Cross Entropy of the the true distributions  $p$  and estimated distribution  $q$ :

$$H(p, q) = E_p[-\log q] = - \sum_{x \in \mathcal{X}} p(x) \log q(x) \quad (17)$$

## 2.4 Gradient Descent

Want to find  $w$ ,  $b$  that minimize  $J(w, b)$ .

Repeat:

$$w := w - \alpha \frac{dJ(w, b)}{dw} \quad (18)$$

$$b := b - \alpha \frac{dJ(w, b)}{db} \quad (19)$$

$\alpha$ : learning rate

## 2.5 Derivatives

$$f(a) = 3a, \frac{df(a)}{da} = 3 = \frac{d}{da} f(a) \quad (20)$$

$$f(a) = a^2, \frac{d}{da} f(a) = 2a \quad (21)$$

$$f(a) = a^3, \frac{d}{da} f(a) = 3a^2 \quad (22)$$

$$f(a) = \log_e a = \ln a, \frac{d}{da} f(a) = \frac{1}{a} \quad (23)$$

$$f(x) = \log_a x, \frac{d}{dx} f(x) = \frac{1}{x \ln a} \quad (24)$$

$$f(x) = a^x, \frac{d}{dx} f(x) = a^x \ln a \quad (25)$$

$$\log_a b = \frac{\log_c b}{\log_c a} = \frac{\ln b}{\ln a} \quad (26)$$

## 2.6 Computation Graph

$$J(a, b, c) = 3(a + bc) \quad (27)$$

$$= 3(a + u) \quad (28)$$

$$= 3v \quad (29)$$

$$u = bc \quad (30)$$

$$v = a + u \quad (31)$$

$$J = 3v \quad (32)$$

## 2.7 Derivatives with a Computation Graph

$$a = 5, b = 3, c = 2$$

$$\frac{dJ}{dv} = 3 \quad (33)$$

$$\frac{dJ}{da} = \frac{dJ}{dv} \frac{dv}{da} \quad (34)$$

$$= 3 * 1 \quad (35)$$

$$= 3 \quad (36)$$

$$\frac{dJ}{du} = \frac{dJ}{dv} \frac{dv}{du} \quad (37)$$

$$= 3 \quad (38)$$

$$\frac{dJ}{db} = \frac{dJ}{du} \frac{du}{db} \quad (39)$$

$$= 3 * c \quad (40)$$

$$= 3 * 2 \quad (41)$$

$$= 6 \quad (42)$$

$$\frac{dJ}{dc} = \frac{dJ}{du} \frac{du}{dc} \quad (43)$$

$$= 3 * b \quad (44)$$

$$= 3 * 3 \quad (45)$$

$$= 9 \quad (46)$$

## 2.8 Logistic Regression Gradient Descent

### 2.8.1 Logistic regression recap

$$z = w^T x + b \quad (47)$$

$$\hat{y} = a = \sigma(z) \quad (48)$$

$$\mathcal{L}(a, y) = -(y \log(a) + (1 - y) \log(1 - a)) \quad (49)$$

### 2.8.2 Logistic regression derivatives

$$z = w_1 x_1 + w_2 x_2 + b \rightarrow a = \sigma(z) \rightarrow \mathcal{L}(a, y)$$

$$da = \frac{d\mathcal{L}(a, y)}{da} \quad (50)$$

$$= -\frac{y}{a} + \frac{1 - y}{1 - a} \quad (51)$$

$$dz = \frac{d\mathcal{L}(a, y)}{dz} \quad (52)$$

$$= \frac{d\mathcal{L}}{da} * \frac{da}{dz} \quad (53)$$

$$= \left(-\frac{y}{a} + \frac{1 - y}{1 - a}\right) * a * (1 - a) \quad (54)$$

$$= -(1 - a)y + a(1 - y) \quad (55)$$

$$= a - y \quad (56)$$

$$\frac{da}{dz} = a * (1 - a) \quad (57)$$

$$= \frac{1}{1 + e^{-z}} * \left(1 - \frac{1}{1 + e^{-z}}\right) \quad (58)$$

$$dw_1 = x_1 * dz \quad (59)$$

$$dw_2 = x_2 * dz \quad (60)$$

$$db = dz \quad (61)$$

Repeat:

$$w_1 := w_1 - \alpha dw_1 \quad (62)$$

$$w_2 := w_2 - \alpha dw_2 \quad (63)$$

$$b := b - \alpha db \quad (64)$$

## 2.9 Gradient Descent on $m$ examples

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(a^{(i)}, y(i)) \quad (65)$$

$$\frac{d}{dw_1} J(w, b) = \frac{1}{m} \sum_{i=1}^m \frac{d}{dw_1} \mathcal{L}(a^{(i)}, y(i)) \quad (66)$$

$$= \frac{1}{m} \sum_{i=1}^m dw_1^{(i)} \quad (67)$$

## 2.10 Vectorization

Vectorized, for GPU.

```
z = np.dot(w, x) + b
```

## 2.11 More Vectorization Examples

### 2.11.1 Neural network programming guideline

Whenever possible, avoid explicit for-loops.

```
u = np.dot(A, x)
```

### 2.11.2 Vectors and matrix valued functions

```
import numpy as np
```

```
u = np.exp(v)
np.log(v)
np.abs(v)
np.maximum(v, 0)
np.pow(v, 2)
np.divide(1, v)
```

## 2.12 Vectorizing Logistic Regression

## 2.13 Vectorizing Logistic Regression's Gradient Computation

## 2.14 Broadcasting in Python

```
np.add((1.0, 2.0), 2.0)    # (3.0, 4.0)
```

## 2.15 Explanation of logistic regression cost function (Optional)

### 2.15.1 Logistic regression cost function

$\hat{y} = \sigma(w^T x + b)$  where  $\sigma(z) = \frac{1}{1+e^{-z}}$

$$\hat{y} = p(y=1|x) \quad (68)$$

If  $y = 1$ :  $p(y|x) = \hat{y}$

If  $y = 0$ :  $p(y|x) = 1 - \hat{y}$

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{(1-y)} \quad (69)$$

$$y = 1 : p(y|x) = \hat{y} \quad (70)$$

$$y = 0 : p(y|x) = 1 - \hat{y} \quad (71)$$

$$\log p(y|x) = y \log \hat{y} + (1 - y) \log 1 - \hat{y} \quad (72)$$

$$= -\mathcal{L}(\hat{y}, y) \quad (73)$$

### 2.15.2 Cost on $m$ examples

Maximum likelihood estimation:

$$\log p(\dots) = \sum_{i=1}^m -\mathcal{L}(\hat{y}, y) \quad (74)$$

$$= -\sum_{i=1}^m \mathcal{L}(\hat{y}, y) \quad (75)$$



Minimize cost:

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}, y) \quad (76)$$